


Introduction to word2vec (skip-gram)

Farid Khafizov
Aug 19, 2021



<http://projector.tensorflow.org/>

Content

1. Introduction
 - a. Define “meaning”
 - b. Learning (new) words
2. Problem statement
3. Skip-gram Algorithm [1,2]
4. Demonstration
5. How technique becomes technology
6. Some Fun Examples

[1] <https://arxiv.org/pdf/1301.3781.pdf> Mikolov 2013.09

[2] <https://arxiv.org/pdf/1310.4546.pdf> Mikolov 2013.10

Problem statement: identify words
with similar meaning

Q1: Can machines learn the meaning of words?

What do we mean by "meaning"?

Winograd schema challenge

- ❑ The city councilmen refused the demonstrators a permit because **they** feared violence.
- ❑ The city councilmen refused the demonstrators a permit because **they** advocated violence.

Natural language understanding (NLU) is central in NLP.

Kevin Gimpel: “the biggest open problems (in NLP) are related to natural language understanding. [...] we should develop systems that read and understand text the way a person does”

How do we learn meaning of words?

1. He handed her a glass of **bardiwac**.
2. Pete staggered to his feet, face flushed from too much **bardiwac**.
3. Malbec, one of the lesser-known **bardiwac** grapes, responds well to Australia's sunshine.
4. The drinks were delicious: blood-red **bardiwac** as well as light, sweet Rhenish.
5. I dined off bread and cheese and this excellent **bardiwac**.
6. Beef dished are made to complement the **bardiwac**.

Answer to Q1

- Machines can not learn “meaning” of words.
We are not there yet.
- However machines can figure out if a word representation is close to another word’s representation in a given context.


Embedding

Text = “Ann was **happy** to be finally home.

Pete was **thrilled** to see Ann returned.”

One-hot encoding (sparse embedding):

- “ann” = $[1, 0, \dots, 0, \dots, 0, \dots, 0] = \underline{v_{\text{ann}}}$
- “was” = $[0, 1, \dots, 0, \dots, 0, \dots, 0]$
- “**happy**” = $[0, 0, \dots, 1, \dots, 0, \dots, 0] = \underline{v_{\text{happy}}}$
- “**thrilled**” = $[0, 0, \dots, 0, \dots, 1, \dots, 0] = \underline{v_{\text{thrilled}}}$
- “returned” = $[0, 0, \dots, 0, \dots, 0, \dots, 1]$



$$\text{Similarity}(v_{\text{happy}}, v_{\text{thrilled}}) = 0$$

Sparse embedding vs Dense embedding

Which one is better?

Dense embedding:

- “ann” = $[0.3, 0.2, 0.3, 0.2]$
- “was” = $[0.1, 0.1, 0.3, 0.5]$
- “**happy**” = $[0.00, 0.21, 0.11, 0.69]$
- “**thrilled**” = $[0.01, 0.20, 0.13, 0.66]$
- “returned” = $[0.4, 0.22, 0.28, 0.1]$



$$\text{Similarity}(v_{\text{happy}}, v_{\text{thrilled}}) = 0.98$$

Illustration

<https://turbomaze.github.io/word2vecjson>

Similar Words

Enter a word and see words with similar vectors.



| | | |
|--------------|--------------------|------------|
| happy | | List words |
| happy | 1 | |
| glad | 0.7408889975979509 | |
| pleased | 0.6632168300400823 | |
| thrilled | 0.651404662202906 | |
| satisfied | 0.6437953159801051 | |
| proud | 0.6360420824282996 | |
| delighted | 0.6272380305363671 | |
| disappointed | 0.6269950290513414 | |
| excited | 0.6247663940787201 | |
| happier | 0.624462516505867 | |

A simpler problem

Separate words according to their context
via embedding words to 2D space.

Note: for us a word embedding is a 2D vector, or
a point on a plane

Random Embedding and Desired Embedding



| word | desired_embedding |
|-------|-----------------------------|
| birds | <code>[-0.24, -0.46]</code> |
| fish | <code>[1.46, -0.05]</code> |
| sky | <code>[-0.3, -0.36]</code> |
| water | <code>[1.29, 0.06]</code> |


Find embedding reflecting semantics of words in a given context.

GIVEN
TEXT:

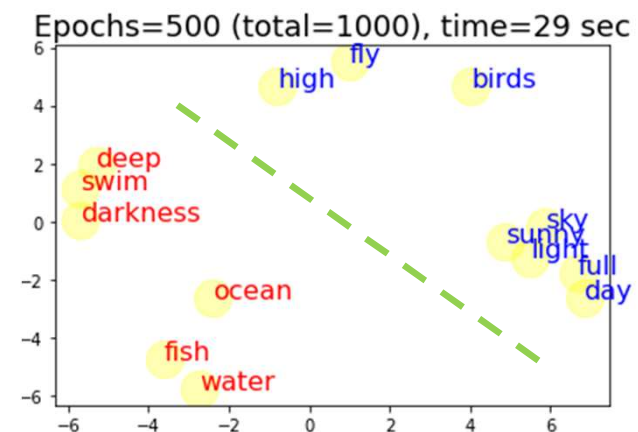
- Fish swim in deep water.
- Ocean is very deep.
- Fish swim in darkness.
- Birds are high in the sky.
- Birds fly very high.
- On a sunny day the sky is full of light.

VOCABULARY (14 words)= ['birds', 'darkness', 'day', 'deep', 'fish', 'fly', 'full', 'high', 'light', 'ocean', 'sky', 'sunny', 'swim', 'water']

and RANDOM embedding:



Find a BETTER embedding:



APPROACH:

- Use context
- WINDOW SIZE = 2
- Embedding Dim=2
- 44 CONTEXT PAIRS

| | input | label |
|----|-------|-------|
| 0 | fish | swim |
| 1 | fish | deep |
| 2 | swim | fish |
| 3 | swim | deep |
| 4 | swim | water |
| 43 | light | full |

Method

Pairs of words for training

| | center | context |
|---|--------|---------|
| 0 | fish | swim |
| 1 | fish | deep |
| 2 | swim | fish |
| 3 | swim | deep |
| 4 | swim | water |

...

| | | |
|----|-------|-------|
| 39 | full | day |
| 40 | full | sky |
| 41 | full | light |
| 42 | light | sky |
| 43 | light | full |

- Fish swim **in** deep water.
- Ocean **is** very deep.
- Fish swim **in** darkness.
- Birds **are** high **in the** sky.
- Birds fly very high.
- On **a** sunny day the sky **is** full **of** light.

Random
initialization
of W_0



| vocab | init_embedding |
|----------|----------------|
| birds | [0.04, 0.15] |
| darkness | [-0.8, -0.68] |
| day | [0.27, 1.54] |
| deep | [1.57, -0.06] |
| fish | [0.61, 1.68] |
| fly | [1.2, 1.08] |
| full | [-1.68, -2.54] |
| high | [-0.83, 0.72] |
| light | [0.69, -0.93] |
| ocean | [-1.51, -3.47] |
| sky | [-0.54, 0.37] |
| sunny | [-0.97, -0.8] |
| swim | [-1.3, -0.15] |
| water | [0.89, 0.31] |

Sketch of the Training Process

For ('birds' , 'fly')

birds: $x \rightarrow x = W_0[0,:] = (0.5, -0.4)$

fly: $y \rightarrow y' = W_1[:,5] = (-0.5, 0.5)^t$

$$W_0 = \begin{bmatrix} 0.5 & -0.4 \\ -0.4 & 0.6 \\ 0.5 & -0.3 \\ 0.5 & -0.1 \\ -0.4 & -0.4 \\ 0.2 & -0.2 \\ -0.2 & -0.2 \\ 0.2 & 0.5 \\ -0.2 & 0.1 \\ 0.1 & -0.1 \\ 0.5 & 0.1 \\ -0.6 & -0.1 \\ 0.3 & -0.2 \\ -0.6 & -0.1 \end{bmatrix}$$

y'

$$W_1 = \begin{bmatrix} -0.2 & 0.5 & 0. & 0.5 & -0.3 & -0.5 & -0.1 & 0.1 & 0.2 & -0.1 & -0.3 & -0.5 & -0.1 & 0.2 \\ -0.3 & -0.2 & -0.6 & -0.6 & 0.6 & 0.5 & 0.3 & 0.3 & 0.6 & 0.1 & -0.1 & 0.1 & -0.1 & 0.1 \end{bmatrix}$$

$x \cdot y'$

$$x \cdot W_1 = \begin{bmatrix} 0.02 & 0.33 & 0.24 & 0.49 & -0.39 & -0.45 & -0.17 & -0.07 & -0.14 & -0.09 & -0.11 & -0.29 & -0.01 & 0.06 \end{bmatrix}$$

$$x = W_0[0,:]$$

$$x \cdot W_1$$

$$\text{dot-product}(s) = x \cdot W_1$$

$$y' = W_1[:,5]$$

$$x \cdot y'$$

Compute Utility Function
(average log likelihood)

Optimize Weights W_0, W_1

Целевая функция

$$P(\overset{\text{'fly'}}{w_{context}} | \overset{\text{'birds'}}{w_{center}}) = P(y|x) = \frac{\exp(x \cdot y')}{\sum_{v=1}^{|V|} \exp(x \cdot y'_v)} = \frac{\exp(x \cdot y')}{D(x, W_1)}$$

$$\log P(y_{t+j}|x_t) = x_t \cdot y'_{t+j} - \log D(x_t, W_1)$$

$$U(W_0, W_1) = \frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log P(y_{t+j}|x_t)$$




Given a sequence of training words $w_1, w_2, w_3, \dots, w_T$, the objective of the Skip-gram model is to maximize $U(W_0, W_1)$.

$$\max_{W_0, W_1} U(W_0, W_1)$$

In our example: $|V|=14$, $T=44$ is the size of training set, $c=2$ is the context window size.

Results

Initial embedding



Q2: Can we simplify skip-gram?

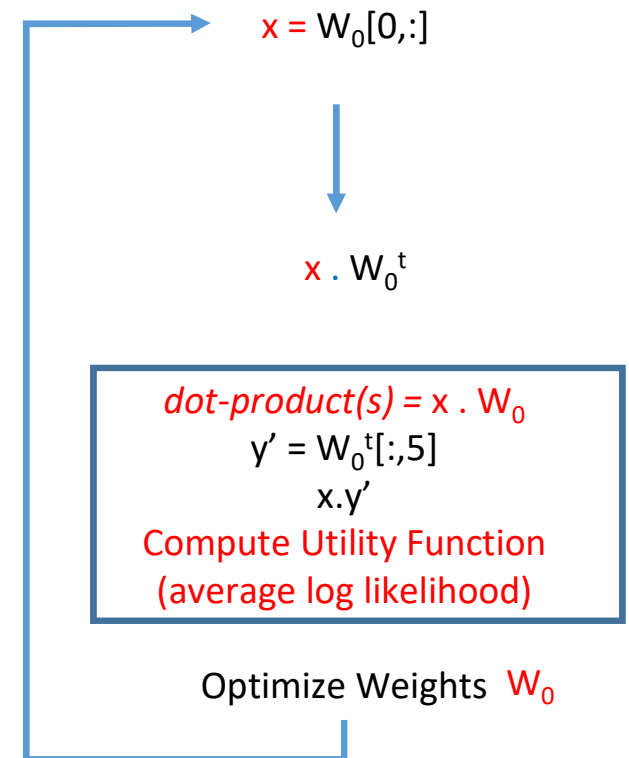
Let $W_1 = W_0^t$

For ('birds' , 'fly')

birds: $x \rightarrow x = W_0[0,:] = (0.5, -0.4)$

fly: $y \rightarrow y' = W_0^t[:,5] = (0.2, -0.2)^t$

x [0.5 -0.4]
[-0.4 0.6]
[0.5 -0.3]
[0.5 -0.1]
[-0.4 -0.4]
 $W_0 =$ [0.2 -0.2] y
[-0.2 -0.2]
[0.2 0.5]
[-0.2 0.1]
[0.1 -0.1]
[0.5 0.1]
[-0.6 -0.]
[0.3 -0.2]
[-0.6 -0.]



Part4: Demo

https://github.com/fkhafizov/w2v_intro

How technique becomes technology

Q3: How to scale?

How many words are out there?

- 470k in Webster's 3rd New International Dictionary
- 170k in current use per Oxford English Dictionary
- 20k-35k used by average native speaker

Technique ==> Technology

1. Distributed architecture and vocab optimization

2. **Subsampling frequent words** ==> Speed up computation

Idea: Discard each training word w_i with probability $P(w_i)$

$$P(w_i) = \begin{cases} \sqrt{1 - \frac{t}{f(w_i)}} & , \quad \text{if } f(w_i) > t \\ 0 & , \quad \text{otherwise} \end{cases}$$

3. **Negative sampling (NEG)** ==> Improves optimization.

Idea: add words from "negative" sample to computation of utility function

Subsampling of frequent words

In very large corpora, the most frequent words can easily occur hundreds of millions of times (e.g., “in”, “the”, and “a”). Such words usually provide less information value than the rare words. For example, while the Skip-gram model benefits from observing the co-occurrences of “France” and “Paris”, it benefits much less from observing the frequent co-occurrences of “France” and “the”, as nearly every word co-occurs frequently within a sentence with “the”. This idea can also be applied in the opposite direction: the vector representations of frequent words do not change significantly after training on several million examples.

To counter the imbalance between the rare and frequent words, we used a simple subsampling approach: each word w_i in the training set is discarded with probability computed by the formula

$$P(w_i) = 1 - \sqrt{\frac{t}{f(w_i)}} \quad (5)$$

where $f(w_i)$ is the frequency of word w_i and t is a chosen threshold, typically around 10^{-5} . We chose this subsampling formula because it aggressively subsamples words whose frequency is greater than t while preserving the ranking of the frequencies. Although this subsampling formula was chosen heuristically, we found it to work well in practice. It accelerates learning and even significantly improves the accuracy of the learned vectors of the rare words, as will be shown in the following sections.

[Mikolov 2013.10]

Negative Sampling

simplify NCE as long as the vector representations retain their quality. We define Negative sampling (NEG) by the objective

$$\log \sigma(v'_{w_O}{}^\top v_{w_I}) + \sum_{i=1}^k \mathbb{E}_{w_i \sim P_n(w)} \left[\log \sigma(-v'_{w_i}{}^\top v_{w_I}) \right] \quad (4)$$

which is used to replace every $\log P(w_O|w_I)$ term in the Skip-gram objective. Thus the task is to distinguish the target word w_O from draws from the noise distribution $P_n(w)$ using logistic regression, where there are k negative samples for each data sample. Our experiments indicate that values of k in the range 5–20 are useful for small training datasets, while for large datasets the k can be as small as 2–5. The main difference between the Negative sampling and NCE is that NCE needs both

[Mikolov 2013.10]

| | | | | |
|--------------------|--------------------|-----------------|------------------------------|---|
| BAD SOLUTION: | $w_I=x='birds'$ | $w_O=y='ocean'$ | $w_n=n='fly'$ | |
| $x = [1.1 \ -3.9]$ | $n = [0.1 \ -4.0]$ | $xn = 15.7,$ | $\text{sig}(-xn) = 1.5e-07,$ | $\log(\text{sig}(-uv)) = -15.7 \quad \ll 0$ |
| $x = [1.1 \ -3.9]$ | $y = [-0.1 \ 3.2]$ | $xy = -12.6,$ | $\text{sig}(xy) = 3.4e-06,$ | $\log(\text{sig}(uv)) = -12.6 \quad \ll 0$ |

| | | | | |
|--------------------|--------------------|---------------|---------------------------|---|
| GOOD SOLUTION: | $w_I=x='birds'$ | $w_O=y='fly'$ | $w_n=n='ocean'$ | |
| $x = [1.1 \ -3.9]$ | $y = [0.1 \ -4.0]$ | $xy = 15.7,$ | $\text{sig}(xy) = 0.99,$ | $\log(\text{sig}(xy)) = -1.5e-07 \quad \sim 0$ |
| $x = [1.1 \ -3.9]$ | $n = [-0.1 \ 3.2]$ | $xn = -12.6,$ | $\text{sig}(-xn) = 0.99,$ | $\log(\text{sig}(-xn)) = -3.4e-06 \quad \sim 0$ |



More details on NEG can be found in section 2 of [Goldberg, 2014.02] <https://arxiv.org/pdf/1402.3722.pdf>

“Word Embedding” search trend



Fun stuff and Interesting observations

$$v(\text{king}) - v(\text{man}) + v(\text{woman}) \approx v(\text{queen})$$




$$\begin{aligned} v(\text{Russia}) + v(\text{river}) &\approx v(\text{Volga River}) \\ v(\text{Germany}) + v(\text{capital}) &\approx v(\text{Berlin}) \end{aligned}$$

Word Algebra

$$\text{France} + (\text{Moscow} - \text{Russia}) = X$$

$$X = ?$$

Geographical Words



Word Embedding Demo

<https://turbomaze.github.io/word2vecjson/>

Similar Words

Enter a word and see words with similar vectors.

| | |
|-------------------------------------|---|
| <input type="text" value="doctor"/> | <input type="button" value="List words"/> |
| doctor | 1.0000000000000002 |
| physician | 0.7806019127031032 |
| doctors | 0.7476568731527384 |
| surgeon | 0.6793393714387082 |
| dentist | 0.6785442117848048 |
| nurse | 0.6319524227288814 |
| psychiatrist | 0.614703850361634 |
| medical | 0.5671389130686404 |
| clinic | 0.5499804910039348 |
| therapist | 0.5283346636619084 |

Word Algebra

Enter all three words, the first two, or the last two and see the words that result.

+ (-) =

spain 0.7905075552539405
madrid 0.7650632609053115


Word Algebra

Enter all three words, the first two, or the last two and see the words that result.

+ (-) =

shanghai 0.779147686225549
china 0.7691120887831315
chinese 0.6720659726186436
hu 0.5964189163973439
yuan 0.5946876191518002

PCA projection of the 1000-dimensional Skip-gram vectors



Mikolov et al, 2013

References

- [0] Mikolov et al, 2013, Linguistic Regularities in Continuous Space Word Representations, <https://www.microsoft.com/en-us/research/wp-content/uploads/2016/02/rvecs.pdf>
- [1] Mikolov et al 2013.09 <https://arxiv.org/pdf/1301.3781.pdf>
- [2] Mikolov et al 2013.10 <https://arxiv.org/pdf/1310.4546.pdf>